

To Decompose or not to Decompose: A Comparison between Different Neural Networks and GBM models for Stock Prices Predictions

Hassan Saadi

Abstract

In this work, 16 stocks are considered to predict the adjusted closing price for the next 40 trading days. The methods that were used to make predictions are: Long Short Term Memory (LSTM), Temporal Convolution Networks (TCN), decomposing the signal and applying different combinations of LSTM and TCN on the decomposed parts, and stochastic process-geometric Brownian motion. The historical data from Yahoo Finance from 2015-2018 was used to build all the models. Finally, the output of each model is compared to the actual adjusted closing price of each stock. The results show that applying a TCN yielded on average the best results and it has the smallest number of parameters among all neural networks models. However, decomposition with a trend, seasonality, and noise components of a time series yielded the best results for some stocks.

1 Introduction

Forecasting the stock prices accurately is a task that many companies are trying to solve and optimize on a daily basis. In the last few decades, the field of machine learning has been developing rapidly and the usage of neural networks is becoming more popular because of their success in modeling complex functions. Some types of neural networks are suited for certain tasks. For example, LSTM and TCN neural networks are good for sequences; hence, they are good for time series predictions; and convolution neural networks

(CNN) are good for images. We also decomposed the time series into three components: trend, seasonality, and residual in order to study the impact of this decomposition on the performance.

Stochastic modeling is an alternative method to model a system, and it's different from a neural network approach. In finance, many people believe that the stock prices follow a geometric Brownian motion (GBM) behavior. That's to say, the returns of stock between two consecutive points in time are normally distributed and independent from each other. In this model the drift and volatility are constant, and are estimated from historical data. This assumption doesn't match in real-world time series, but future work can incorporate time dependent estimations for these parameters.

In this work, we didn't consider autoregressive models like ARIMA because the authors of this work [1] showed the superiority of LSTMs in predicting the adjusted closing prices where the dataset was large. Moreover, the prediction period in our work is 40 trading days which we considered to be not a small window of predictions, and we have 1000 training points which is not a small dataset, so a neural network can have a chance to learn it. ARIMA can be a good model if the time series is stationary or can be made stationary. But if the time series cannot be made stationary trivially, then using ARIMA models becomes not ideal. Stationarity is not required for neural networks because they can learn non-linear relationships.

To the the best of our knowledge, the literature does not contain a comparison study between LSTM, TCN, decomposed time series with LSTM and TCN, and GBM. We decided to perform these experiments on 16 stocks, and predict the adjusted closing price for the next 40 trading days in order to assess the performance of TCN and the decomposition idea. In this work, section 2 has a short literature review and related work. In section 3, has the methodology. In section 4 is the experimental setup. In section 5, the results of all the models. And finally, in section 6 conclusion and future work.

2 Short literature review and related work

In [2] work, they compared four types of neural networks (RNN,LSTM,CNN, and MLP) to predict two stock markets price. In [3] work, they also compared

LSTM, CNN, MLP, and one attention-based neural network to predict the next day's index price from three different financial markets. The attention-based model performed the best in their experiments. Another group created a new neural architecture to deal with the temporal structure that a time series has by inventing Temporal Convolution Network (TCN) [4].

3 Methodology

3.1 LSTM

Our first model is an LSTM. The formulation of LSTM is as follows

$$\mathbf{i}_t = \sigma_g(\mathbf{W}_i \mathbf{x}_t + \mathbf{U}_i \mathbf{h}_{t-1} + \mathbf{b}_i) \quad (1)$$

$$\mathbf{f}_t = \sigma_g(\mathbf{W}_f \mathbf{x}_t + \mathbf{U}_f \mathbf{h}_{t-1} + \mathbf{b}_f) \quad (2)$$

$$\tilde{\mathbf{c}}_t = \sigma_c(\mathbf{W}_c \mathbf{x}_t + \mathbf{U}_c \mathbf{h}_{t-1} + \mathbf{b}_c) \quad (3)$$

$$\mathbf{c}_t = \mathbf{f}_t \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t \quad (4)$$

$$\mathbf{o}_t = \sigma_g(\mathbf{W}_o \mathbf{x}_t + \mathbf{U}_o \mathbf{h}_{t-1} + \mathbf{b}_o) \quad (5)$$

$$\mathbf{h}_t = LSTM(\mathbf{x}_t, \mathbf{h}_{t-1}) = \mathbf{o}_t \odot \sigma_h(\mathbf{c}_t) \quad (6)$$

where \odot means element-wise multiplication, σ_g is the logistic sigmoid function, σ_c is the hyperbolic tangent function, and σ_h is the hyperbolic tangent function which are used as recurrent activation functions. LSTM architecture has memory cells to store and output information. This property aids a recurrent neural network to find long-range temporal dependencies. \mathbf{i} , \mathbf{f} , \mathbf{c} , and \mathbf{o} are the input gate, forget gate, cell activation, and output gate respectively.

3.2 TCN

TCN has two main ideas: The convolutions are causal, i.e., no information is leaked from future to past; and the architecture can have a sequence of any length as input and map it to an output sequence that has the same length. The first idea is accomplished by making convolutions where an output at time t is convolved only with elements from time t and in the past in the previous layer. The second idea is accomplished by utilizing a $1D$ fully convolutional network architecture such that each hidden layer has the same size as the input layer. Dilated convolutions are also used here to enable a

larger receptive field to capture dependencies at larger scales. For an input sequence $\mathbf{x} \in \mathbb{R}^n$ and a filter $h : \{0, \dots, k-1\} \rightarrow \mathbb{R}$, the dilated convolutional operation F on element s on the sequence is

$$F(s) = (\mathbf{x} *_d h)(s) = \sum_{i=0}^{k-1} h(i) \cdot \mathbf{x}_{s-d \cdot i} \quad (7)$$

where d is the dilation factor, k is the size of the filter, and $s - d \cdot i$ captures the direction in the past.




Figure 1: A dilated Causal convolution with dilation factors $d = 1, 2, 4$ and filter size $k = 3$ [5].

3.3 Decomposition

A time series can be decomposed into three parts: trend, seasonality, and residual. In this work, the "*stat_decompose*" library in python was used in order to the decomposition. Note that, the trend is obtained by using convolution filter as a moving average of the last 10 days. After decomposition, I applied different combinations of LSTM(s) and TCN(s) to these components. In total, there is 8 models that incorporate decomposition: (L,L,L), (L,L,T), (L,T,L), (L,T,T), (T,L,L), (T,L,T), (T,T,L), and (T,T,T); where L and T mean LSTM and TCN were used respectively; and the first, second, and third entries in each model correspond to trend, seasonality, and residual part of the time series respectively. So for example, (L,L,T) model means

that two LSTMs are used to model trend and seasonality separately and one TCN is use to model the residual.

3.4 GBM

A Wiener process (Brownian Motion) is a Markov stochastic process with zero drift and variance rate of one; $dW = \epsilon\sqrt{dt}$ where ϵ is sampled from a normal distribution. It has the following properties:

1. $W_0 = 0$
2. $W_{s+t} - W_t \in \mathcal{N}(0, s)$ for $0 \leq s \leq t \leq T$
3. $\{W_t\}_{t \geq 0}$ has stationary and independent non-overlapping increments.

A generalized Wiener process has a drift term and has the form $dX(t) = \mu dt + \sigma dW(t)$. If the stochastic process is defined as $X(t) = \log(S(t))$ because the prices are assumed to follow a lognormal distribution since prices can't be negative, then we obtain the geometric Brownian motion as

$$\frac{dS(t)}{S(t)} = \mu dt + \sigma dW(t) \quad (8)$$

This can be solved for $t > 0$ by integrating both sides as

$$S(t) = S(0) + \mu \int_0^t S(\tau) d\tau + \sigma \int_0^t S(\tau) dW(\tau) \quad (9)$$

By using Ito's formula we obtain

$$d \log(S(t)) = \left(\mu - \frac{\sigma^2}{2}\right) dt + \sigma dW(t) \quad (10)$$

Hence, after integrating both sides

$$S(t) = S(0)e^{(\mu - \frac{1}{2}\sigma^2)t + \sigma W(t)} = S(0)e^{X(t)} \quad (11)$$

Since we are working with discrete data, we can make progressive predictions from time t_0 incrementally as $t_0 < t_1 < t_2 < \dots < t_N$ by generating N iid $\mathcal{N}(0, 1)$ random variables z_1, \dots, z_N . Note that this can be done because for any $0 \leq s < t$ we have $S(t) = S(0) \frac{S(s)}{S(0)} \frac{S(t)}{S(s)} = S(0)e^{X(s)} \times e^{X(s)-X(t)}$, and

we already know that $X(s)$ is independent of the increment $X(t) - X(s)$; hence, $\frac{S(s)}{S(0)}$ and $\frac{S(t)}{S(s)}$ are independent lognormals. Generalizing the idea, let $Y_i = \frac{S_{t_i}}{S_{(t_{i-1})}}$, we have $S(t_k) = S(t_{k-1})Y_k = S_0Y_1Y_2...Y_k$, with $Y_i = e^{\sigma z_i + \mu}$ because $t_i - t_{i-1} = 1$ in our case because we are interested in the daily adjusted price of a stock. Therefore, we obtain

$$\begin{aligned} S_{t_k} &= S_0 \prod_{i=1}^k e^{\mu - \frac{1}{2}\sigma^2 + \sigma z_i} \\ &= S_0 e^{(\mu - \frac{1}{2}\sigma^2)k + \sigma \sum_{i=1}^k z_i} \\ &= S_0 e^{(\mu - \frac{1}{2}\sigma^2)t_k + \sigma W_k} \end{aligned} \tag{12}$$

where we replaced k with t_k , and $W_k = \sum_{i=1}^k z_i$.

4 Experimental Section

4.1 Datasets

All experiments are performed on 16 stocks from the S&P 500 list: Apple (AAPL), IBM (IBM), Tesla Inc. (TSLA), Microsoft Corp. (MSFT), Facebook (FB), Google Inc. (GOOGL), Procter & Gamble (PG), JPMorgan Chase & Co. (JPM), Netflix Inc. (NFLX), Intel Corp. (INTC), Adobe Inc. (ADBE), Johnson & Johnson (JNJ), Goldman Sachs Group (GS), Morgan Stanley (MS), (NDAQ), General Motors (GM). The historical data was taken from Yahoo Finance from 2015-2018, and the goal is to predict the adjusted close price for the next 40 trading days; i.e. the months of January and February of 2019.

4.2 Evaluation Metrics

Two metrics are adopted in this study, $RMSE$ and MAE :

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (x_i - \hat{x}_i)^2}{N}} \tag{13}$$

$$MAE = \frac{\sum_{i=1}^N |x_i - \hat{x}_i|}{N} \tag{14}$$

where x_i is the true value, \hat{x}_i is the predicted value, N is the total number of test data points.

4.3 Hyperparameters

If only one LSTM is used for all stocks, then its properties are: One layer, 300 neurons, 100 epochs, 64 batch size, time step of 10, mean squared error loss, and adam optimizer. The number of trainable parameters is 362701.

In all experiments, the TCNs properties are: filter size of 2, 100 epochs, time step 10, dilations of $\{1, 2, 4, 8\}$, mean squared error loss, and adam optimizer. The number of trainable parameters is 87937.

If an LSTM is used on the decomposed components then its properties are: One layer, 100 neurons, 100 epochs, 64 batch size, time step of 10, mean squared error loss, and adam optimizer. The number of trainable parameters is 40901. Note that the number of neurons in this LSTM is one-third of the LSTM described earlier.

4.4 Machine Specs

All experiments were run on a laptop with 1.6 GHz Dual-Core Intel Core i5 and 8GB of main memory.

5 Results

In table (1), the average RMSE and MAE over all 16 stocks, and the number of stocks that each model predicted the best according to RMSE and MAE are listed. The best model is shown in bold and colored with blue.

Model	<i>Number of Parameters</i>	<i>Avg. RMSE</i>	<i>Avg. MAE</i>	<i>number of best RMSE</i>	<i>number of best MAE</i>
LSTM	362701	4.22	3.425	8	6
TCN	87937	3.86	3.07	8	9
GBM	0	13.06	11.80	0	0
(L,L,L)	122703	6.07	4.994	0	0
(L,L,T)	169739	6.197	5.076	0	0
(L,T,L)	169739	6.06	4.992	0	0
(L,T,T)	216775	6.195	5.074	0	0
(T,L,L)	169739	4.877	3.88	0	0
(T,L,T)	216775	5.496	4.332	0	0
(T,T,L)	216775	4.875	3.87	0	0
(T,T,T)	263811	5.495	4.33	0	1

Table 1: Average RMSE and MAE for 18 stocks

The best models are TCN, LSTM, (T,T,T) in order. Below are boxplots for the RMSE and MAE values of each model where the mean of each model is shown as a green full circle. The boxplots are without outliers since we only have 16 stocks, and they don't exhibit normality as we can see from the boxplots and after performing QQ plots, Kolmogorov-Smirnov, and Shapiro tests. Hence, the whiskers in the boxplots represent the minimum and maximum values of the dataset (model) range and not the 1.5 IQR.




Figure 2: RMSE values for each model




Figure 3: MAE values for each model

The heat map of RMSE for each stock is listed. The GBM model in this heat map was removed because of its high RMSE and MAE values. In the heat maps, green means low values of z-score, and red means high values of z-score. The tables (2) and (3) of the RMSE and MAE for each stock is listed at the appendix where the best model is shown in bold and colored with blue.

		RMSE values with Z-score colors per row									
Stock	AAPL	0.993	1.786	1.400	1.656	1.400	1.656	1.700	1.737	1.700	1.737
	IBM	1.989	2.983	2.812	3.202	2.812	3.202	3.709	4.143	3.708	4.143
	TSLA	2.613	4.385	3.040	3.358	3.042	3.357	3.127	3.461	3.130	3.460
	MSFT	1.847	1.861	3.296	3.561	3.297	3.561	3.184	3.476	3.185	3.477
	FB	4.542	3.783	5.153	5.347	5.154	5.348	5.250	5.117	5.251	5.118
	GOOGL	19.769	17.223	31.203	26.589	31.191	26.584	21.631	25.437	21.635	25.451
	PG	1.168	3.333	1.677	1.826	1.666	1.814	2.476	2.680	2.460	2.665
	JPM	1.204	1.462	1.451	1.821	1.451	1.821	1.262	1.208	1.261	1.208
	NFLX	10.911	9.591	23.037	22.450	23.036	22.448	14.933	14.777	14.932	14.776
	INTC	1.092	0.950	1.352	1.675	1.352	1.675	1.062	1.182	1.062	1.181
	ADBE	13.612	4.700	10.526	12.297	10.511	12.283	10.040	11.868	10.025	11.853
	JNJ	1.508	1.153	2.952	4.094	2.946	4.090	1.327	2.621	1.324	2.620
	GS	4.000	3.588	5.591	7.525	5.591	7.524	5.112	6.863	5.112	6.861
	MS	0.690	1.205	0.963	0.986	0.963	0.987	1.039	1.090	1.039	1.090
	NDAQ	0.956	3.328	1.786	1.962	1.786	1.962	1.147	1.385	1.147	1.385
	GM	0.629	0.573	0.893	0.811	0.894	0.811	1.042	0.901	1.044	0.902
		LSTM	TCN	LLL	LLT	LTL	LTT	TLL	TLT	TTL	TTT
		Model									

Figure 4: RMSE values with Z-score colors per row for 16 stocks and all models excluding GBM.

MAE values with Z-score colors per row

AAPL	0.670	1.666	1.062	1.263	1.063	1.263	1.412	1.359	1.412	1.359
IBM	1.386	1.975	2.149	2.559	2.148	2.558	3.091	3.613	3.089	3.611
TSLA	1.755	3.795	2.229	2.400	2.232	2.404	2.315	2.514	2.318	2.516
MSFT	1.568	1.535	2.875	3.074	2.876	3.075	2.758	2.998	2.759	2.998
FB	3.465	2.525	3.704	4.098	3.705	4.099	3.443	3.470	3.445	3.474
GOOGL	16.323	13.445	26.196	22.455	26.193	22.452	17.844	20.317	17.842	20.316
PG	0.960	2.952	1.267	1.390	1.257	1.377	2.136	2.360	2.121	2.343
JPM	0.954	1.256	1.168	1.461	1.168	1.461	0.998	0.945	0.998	0.945
NFLX	7.988	7.326	18.732	17.714	18.730	17.712	10.828	10.641	10.827	10.641
INTC	0.789	0.734	1.149	1.425	1.149	1.425	0.841	0.986	0.841	0.986
ADBE	13.140	3.811	10.004	11.535	9.988	11.519	9.416	11.013	9.400	10.997
JNJ	1.219	0.945	2.715	3.445	2.709	3.439	1.050	1.583	1.048	1.581
GS	2.761	2.500	3.720	5.461	3.720	5.460	3.414	4.958	3.414	4.956
MS	0.566	1.055	0.760	0.774	0.759	0.773	0.851	0.901	0.851	0.900
NDAQ	0.783	3.197	1.546	1.596	1.546	1.596	0.885	0.971	0.885	0.971
GM	0.481	0.411	0.632	0.582	0.632	0.582	0.811	0.685	0.813	0.687
	LSTM	TCN	LLL	LLT	LTL	UTT	TLL	TLT	TTL	TTT

Model

Figure 5: MAE values with Z-score colors per row for 16 stocks and all models excluding GBM.

Moreover, we evaluated and plotted the logarithm of the standard deviation on the x-axis and the Hurst exponent on the y-axis of our training set for each stock. In plots (6) and (7), we notice that for high values of standard deviation and Hurst exponent (after around 0.4), TCN performs the best. The colors in these plots denote the best neural network that performed on the written stock according to RMSE or MAE.




Figure 6: $\log(std)$ vs Hurst Exponent, best neural networks performance according to RMSE




Figure 7: $\log(std)$ vs Hurst Exponent, best neural networks performance according to MAE

6 Conclusion and Future Work

In this work, we compared multiple neural network architectures with and without decomposition and a simple GBM model. TCN and LSTM have the best results overall in a decreasing order on 16 stocks to predict the adjusted closing price for the next 40 trading days. According to our results, additive decomposition did not give better results. Furthermore, stocks that have high standard deviations and Hurst exponent were shown to be best predicted by TCN. As future work, one can investigate different methods to improve the performance of these results and study their generalization capabilities. This can be done through building ensembles of these models and using other ways of decomposing the time series such as singular spectrum analysis or STL. Moreover, it would be informative to investigate the relationship between time series properties and neural network architectures and decomposition methods.

References

- [1] S. Siامي-Namini, N. Tavakoli, and A. S. Namin, “A comparison of arima and lstm in forecasting time series,” in *2018 17th IEEE International Conference on Machine Learning and Applications (ICMLA)*. IEEE, 2018, pp. 1394–1401.
- [2] M. Hiransha, E. A. Gopalakrishnan, V. K. Menon, and K. Soman, “Nse stock market prediction using deep-learning models,” *Procedia computer science*, vol. 132, pp. 1351–1362, 2018.
- [3] P. Gao, R. Zhang, and X. Yang, “The application of stock index price prediction with neural network,” *Mathematical and Computational Applications*, vol. 25, no. 3, p. 53, 2020.
- [4] C. Lea, M. D. Flynn, R. Vidal, A. Reiter, and G. D. Hager, “Temporal convolutional networks for action segmentation and detection,” in *proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2017, pp. 156–165.
- [5] S. Bai, J. Z. Kolter, and V. Koltun, “An empirical evaluation of generic convolutional and recurrent networks for sequence modeling,” *arXiv preprint arXiv:1803.01271*, 2018.

7 Appendix

Model Stock	LSTM	TCN	GBM	(L,L,L)	(L,L,T)	(L,T,L)	(L,T,T)	(T,L,L)	(T,L,T)	(T,T,L)	(T,T,T)
AAPL	0.9933	1.7858	2.3101	1.3997	1.6558	1.4000	1.6563	1.7001	1.7368	1.6999	1.7369
IBM	1.9889	2.9826	16.400	2.8124	3.2024	2.8116	3.2015	3.7092	4.1434	3.7082	4.1425
TSLA	2.6129	4.3848	3.7576	3.0401	3.3577	3.0423	3.3572	3.1273	3.4606	3.1295	3.4602
MSFT	1.8471	1.8605	3.5986	3.2961	3.5605	3.2968	3.5613	3.1840	3.4763	3.1847	3.4770
FB	4.5421	3.7830	19.818	5.1527	5.3467	5.1539	5.3483	5.2501	5.1165	5.2510	5.1180
GOOGL	19.769	17.223	34.764	31.202	26.588	31.190	26.584	21.631	25.437	21.634	25.450
PG	1.1678	3.3327	5.2331	1.6773	1.8263	1.6659	1.8136	2.4756	2.6803	2.4601	2.6645
JPM	1.2042	1.4624	3.2384	1.4512	1.8213	1.4512	1.8214	1.2615	1.2079	1.2613	1.2078
NFLX	10.910	9.5910	63.185	23.037	22.449	23.035	22.448	14.932	14.776	14.931	14.775
INTC	1.0915	0.9500	2.5388	1.3520	1.6753	1.3518	1.6751	1.0620	1.1815	1.0620	1.1814
ADBE	13.611	4.6999	18.950	10.525	12.297	10.511	12.282	10.040	11.867	10.025	11.853
JNJ	1.5080	1.1533	3.9341	2.9519	4.0939	2.9455	4.0896	1.3270	2.6209	1.3237	2.6195
GS	4.0003	3.5884	21.124	5.5913	7.5253	5.5910	7.5236	5.1123	6.8625	5.1123	6.8610
MS	0.6903	1.2050	1.6694	0.9629	0.9864	0.9633	0.9867	1.0385	1.0901	1.0386	1.0901
NDAQ	0.9562	3.3277	4.1382	1.7860	1.9624	1.7859	1.9624	1.1466	1.3847	1.1465	1.3847
GM	0.6294	0.5731	4.3646	0.8926	0.8105	0.8936	0.8109	1.0420	0.9008	1.0436	0.9021

Table 2: RMSE for 16 stocks

Model Stock	LSTM	TCN	GBM	(L,L,L)	(L,L,T)	(L,T,L)	(L,T,T)	(T,L,L)	(T,L,T)	(T,T,L)	(T,T,T)
AAPL	0.6699	1.6659	2.0978	1.0624	1.2630	1.0632	1.2632	1.4122	1.3590	1.4121	1.3592
IBM	1.3855	1.9753	14.509	2.1485	2.5592	2.1475	2.5580	3.0905	3.6125	3.0892	3.6112
TSLA	1.7547	3.7946	3.1476	2.2285	2.4004	2.2322	2.4035	2.3146	2.5143	2.3184	2.5162
MSFT	1.5682	1.5354	2.9642	2.8750	3.0740	2.8755	3.0748	2.7583	2.9976	2.7589	2.9984
FB	3.4653	2.5247	17.240	3.7039	4.0976	3.7053	4.0992	3.4428	3.4698	3.4452	3.4740
GOOGL	16.323	13.444	30.560	26.195	22.455	26.192	22.451	17.843	20.317	17.842	20.315
PG	0.9603	2.9516	4.1822	1.2673	1.3901	1.2569	1.3771	2.1359	2.3597	2.1209	2.3430
JPM	0.9535	1.2559	2.9672	1.1676	1.4608	1.1676	1.4610	0.9980	0.9447	0.9978	0.9445
NFLX	7.9876	7.3259	60.671	18.731	17.714	18.729	17.712	10.827	10.641	10.827	10.640
INTC	0.7887	0.7339	1.9854	1.1494	1.4253	1.1491	1.4251	0.8412	0.9858	0.8412	0.9856
ADBE	13.139	3.8107	17.305	10.004	11.534	9.9878	11.518	9.4160	11.013	9.3998	10.997
JNJ	1.2191	0.9451	3.1015	2.7154	3.4449	2.7093	3.4390	1.0502	1.5830	1.0476	1.5806
GS	2.7605	2.4998	19.063	3.7199	5.4605	3.7198	5.4598	3.4141	4.9577	3.4140	4.9561
MS	0.5662	1.0549	1.5041	0.7596	0.7735	0.7594	0.7734	0.8508	0.9006	0.8508	0.9004
NDAQ	0.7829	3.1965	3.5573	1.5457	1.5957	1.5456	1.5956	0.8850	0.9705	0.8850	0.9705
GM	0.4808	0.4110	4.0517	0.6316	0.5819	0.6324	0.5819	0.8112	0.6852	0.8128	0.6867

Table 3: MAE for 16 stocks